ARI Founder Handbook

Foundations, Taste, Execution, and Architecture

March 12, 2026

Contents

Purpose of This Handbook

This handbook is written for a serious builder who wants to train judgment, not just accumulate information.

Its governing idea is simple:

- purpose decides what is worth dedicating your life to
- will decides whether you can endure difficulty
- taste decides what is real, deep, reusable, and worthy

Purpose without taste becomes fanaticism. Will without taste becomes waste. Taste is what lets you choose the right mountain.

What Taste Means

In this context, taste is not preference. It is the ability to distinguish:

- signal from noise
- mechanism from metaphor
- durable leverage from fashionable novelty
- trustworthy systems from impressive demos
- work that matters from work that merely performs importance

The five recurring questions are:

1. Does this touch reality, or only benchmark theater?
2. Does this compound, or reset every hype cycle?
3. Does this unlock reusable capability, or only a narrow artifact?
4. Does this deepen our model of mind, biology, society, or infrastructure?
5. If this succeeds, does it actually matter?

Knowledge System to Rebuild

If starting from a high-school baseline, the best order is:

1. mathematics as the language of structure and change
2. physics and dynamical systems
3. computer science as controlled complexity
4. neuroscience and computation
5. embodiment and robotics
6. biology and augmentation
7. institutions, privacy, and inequality

The unifying lesson is that serious systems are constrained, stateful, adaptive, and embedded in society.

Twelve-Week Operational Plan

How to use the plan

Each week should include:

1. study
2. implementation
3. writing
4. reflection

Each week should also produce one judgment artifact: a short note on what became more real or less convincing.

Week 1 to Week 4

Week 1: Functions, variables, and modeling

Move from school math to causal description. Build tiny simulations of linear growth, exponential decay, and threshold switches. Write one page on why a model is a compressed causal story.

Week 2: Derivatives, accumulation, and state

Learn to think in local change and accumulation through time. Implement Euler integration and observe the effect of step size. Write one page on why discretization already trains engineering taste.

Week 3: Vectors, matrices, and noisy observations

Understand structure and uncertainty together. Implement a matrix-based recurrent update and add sensor noise. Write one page on why real systems are structured and noisy.

Week 4: LIF from first principles

Implement a leaky integrate-and-fire neuron from scratch. Sweep time constant and threshold. Write one page on what LIF preserves and what it discards.

Week 5 to Week 8

Week 5: Synapses, spike trains, and coding

Add synaptic filtering and compare fast and slow decay. Build a two-neuron excitatory or inhibitory motif. Write one page on why not all spikes mean the same thing.

Week 6: Recurrent circuits and competition

Study stabilization, competition, and runaway excitation. Build a small recurrent SNN and identify the difference between strong dynamics and parameter abuse. Write one page on why recurrence is power and danger at once.

Week 7: Embodiment and closed loops

Run the embodied prototype and inspect sensory history, action history, and trajectory. Change one body parameter and observe the shift in behavior. Write one page on why bodies compute too.

Week 8: Connectomes and structure-constrained models

Redraw the prototype network into sensory, interneuron, descending, and motor groups. Justify every connection in words. Write one page on why a wiring diagram is not yet a mind.

Week 9 to Week 12

Week 9: Plasticity and local learning

Implement one toy plasticity rule and track how synaptic weights evolve. Write one page on what information a learning rule actually needs.

Week 10: Interfaces and hidden arbitrariness

Inspect the prototype's sensory mapping and action decoder. Identify which parts are principled and which are hand-tuned. Write one page on why interfaces are where truth leaks.

Week 11: Trust, privacy, and human stakes

Design a toy privacy-first memory system for a personal AI companion. Identify abuse paths. Write one page on why trust is an architecture property before it is a policy statement.

Week 12: Synthesis and founder-level clarity

Write one-page architecture memos for Alive, Realize, and Innocence, plus one page on the shared reusable stack. Write one page on what is worth building.

Non-negotiable rules

1. Never let reading replace implementation.
2. Never let implementation replace writing.
3. Never let ambition excuse vagueness.
4. Record confusions explicitly.
5. Refuse at least one seductive but shallow idea every two weeks.

ARI Taste Standards

Any major direction should pass six tests:

1. reality test
2. mechanism test
3. reusability test
4. trust test
5. dignity test

6. durability test

If a direction fails several of these, it should not receive deep commitment.

Alive

Human problem

Loneliness is not only absence of contact. It is absence of witnessed identity, continuity, and trusted presence.

Irreducible technical object

Not a chatbot, but a persistent relational agent with memory, boundaries, continuity, and possibly embodiment.

Architecture layers

1. identity layer with stable persona and versioned self-model
2. memory layer with episodic, preference, and relational memory
3. interaction layer across text, voice, visual, and possibly embodied channels
4. user sovereignty layer with real memory editing, deletion, and export
5. trust and safety layer with crisis-sensitive behavior and clear limits

Taste standard

Prefer continuity over novelty, genuine care over retention tricks, and agency over dependency.

Realize

Human problem

Biological limits appear as disability, decline, suffering, and hard capability ceilings.

Irreducible technical object

Not an enhancement slogan, but a closed-loop augmentation or restoration system grounded in measurable biology.

Architecture layers

1. sensing layer for noisy biological signal acquisition
2. interpretation layer with uncertainty-aware decoding and subject adaptation
3. intervention layer under strict safety envelopes
4. clinical and validation layer with evidence tracking and auditability
5. human factors layer for comfort, adherence, and informed use

Taste standard

Prefer mechanism over spectacle, restoration over theater, and translational seriousness over futurist language.

Innocence

Human problem

Systemic inequality persists when infrastructure is opaque, extractive, and easy to weaponize.

Irreducible technical object

Not a fairness dashboard, but a privacy-first coordination and service infrastructure that preserves agency under power asymmetry.

Architecture layers

1. identity and consent layer
2. data boundary layer with encryption and access control
3. service orchestration layer across institutions
4. user interface layer that remains understandable and accessible
5. governance resilience layer against abuse and capture

Taste standard

Prefer privacy by architecture over policy theater, inclusion as actual usability over rhetoric, and hard trust boundaries over benevolent assumptions.

Shared ARI Core

The ecosystem should compound through common technical layers:

1. identity and continuity
2. privacy and consent
3. multimodal interaction
4. personal memory and retrieval
5. edge and device integration
6. safety, auditability, and policy enforcement
7. adaptation under user control

The strategic principle is simple: products differ at the surface, but trust, memory, and control architecture should compound underneath.

Final Standard

Choose work that is simultaneously:

- real
- deep
- reusable

- worthy

Alive should make companionship more real, not more addictive.

Realize should make augmentation more grounded, not more theatrical.

Innocence should make infrastructure more trustworthy, not more rhetorical.

If a technical choice strengthens those standards, it is probably aligned. If it weakens them, it is probably a distraction even if it is fashionable.